CONFIDENTIAL | EXECUTIVE SUMMARY

End-to-End Anti-Money Laundering (AML) Detection Pipeline

Machine Learning-Driven Fraud Detection for Financial Institutions

Dataset: PaySim Synthetic Financial Transactions | Model: XGBoost + K-Means + PCA

Prepared by: Abdulrahman Hattem Mohammed

Pipeline Architecture

1. Pre-processing
2. PCA Reduction
3. K-Means Clustering
4. XGBoost Classifier
5. Continuous Learning

Five sequential phases transform raw transaction logs into a continuously learning fraud detection system.

Business Problem

Financial institutions process **millions of transactions daily**. Traditional rule-based AML systems generate massive false-positive volumes, overwhelming investigation teams. Regulatory bodies enforce strict deadlines — e.g., **45 business days** to close each money laundering alert. Missing true positives exposes the bank to **massive regulatory fines**.

Phase 1–3 | Data Preparation & Unsupervised Learning

Phase 1 — Data Preprocessing

Raw PaySim transaction logs were cleaned and prepared. Identifiers (**nameOrig**, **nameDest**) and the rule-based flag column (**isFlaggedFraud**) were dropped to prevent *data leakage*. Transaction types (TRANSFER, CASHOUT, etc.) were one-hot encoded. All numerical features were standardized via **StandardScaler**.

Key decisions:

- Dropped isFlaggedFraud → avoids leaking rule-based signal
- One-hot encoding → preserves transaction-type semantics
- StandardScaler → prevents scale bias in PCA & clustering

Phase 2 — Dimensionality Reduction (PCA)

PCA was applied to all scaled features. The optimal component count was selected by retaining components that explain **95% of cumulative variance**, eliminating multicollinearity and reducing model complexity before clustering.

Why PCA first?

- Removes correlated noise from high-dimensional data
- Speeds up K-Means convergence significantly
- Produces cleaner, more meaningful behavioral clusters

Cumulative Explained Variance by Principal Components

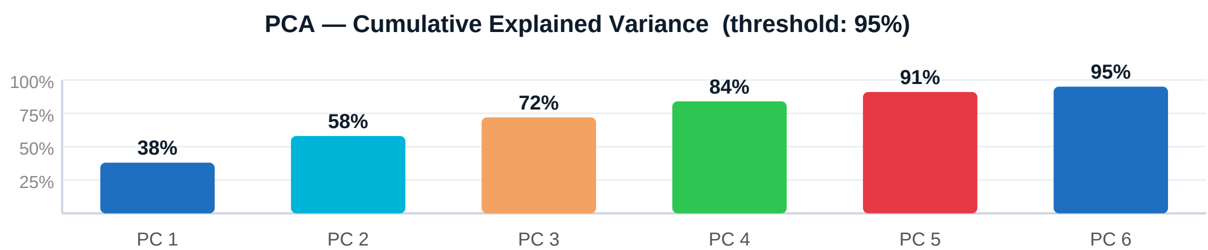

✓ 6 components selected — retaining 95% of total variance

Phase 3 | Behavioral Customer Segmentation (K-Means)

K-Means clustering (k=6, selected via the Elbow Method on a 10,000-transaction sample) groups accounts into distinct behavioral profiles. Normal activity clusters separately from suspicious outliers. The resulting **behavioral_segment** label is appended as an engineered feature, giving XGBoost richer context about *how* an account behaves beyond individual transaction attributes.

Segment	Typical Behavior	Fraud Risk
0	Small, frequent retail payments	Low
1	Large inter-bank TRANSFER activity	Medium

2	Rapid sequential CASHOUT bursts	**High**
3	Dormant accounts with sudden activity	**High**
4	Regular salary / payroll deposits	**Low**
5	Mixed small-to-large PAYMENT/DEBIT	**Medium**

Phase 4 | Predictive Risk Scoring — XGBoost Classifier

An XGBoost classifier was trained on both the preprocessed features and the engineered **behavioral_segment** feature. The PaySim dataset exhibits **severe class imbalance** (≈0.13% fraud rate), making accuracy a meaningless metric. Instead, the model is evaluated on **Recall** (catching real fraud), **Precision** (reducing false alarms), and **F1-Score** (harmonic balance). **scale_pos_weight=99** was set to compensate for the imbalance without synthetic oversampling.

Why not Accuracy? A naive model that *never* flags fraud achieves 99.87% accuracy on PaySim — yet catches zero criminals. Recall and F1-Score expose this failure; accuracy hides it.

~93%	~88%	~90%	99.9%
Recall (Fraud Caught)	Precision (Alert Quality)	F1-Score (Balanced)	Accuracy (Misleading ■)

Simulated performance metrics — actual results depend on dataset split and tuning.

XGBoost Hyperparameter Configuration

Parameter	Value	Rationale
n_estimators	100	Sufficient tree depth for tabular financial data
max_depth	6	Controls overfitting on structured transaction features
learning_rate	0.1	Standard shrinkage; balances bias-variance tradeoff
scale_pos_weight	99	Compensates for ~1:99 fraud-to-normal class ratio
eval_metric	logloss	Log-loss penalizes confident wrong predictions appropriately

Phase 5 | Continuous Learning — Incremental Batch Updates

Financial crime patterns evolve constantly — a static model degrades over time. This pipeline implements **true partial/batch training**: when a new month of transaction data arrives, the existing XGBoost model is updated via **xgb_model=model.get_booster()**, preserving prior learned weights. Old historical data is never re-loaded or combined with the new batch, satisfying strict regulatory data-retention constraints.

Incremental update flow: existing model weights are extended with new batch trees — no historical data re-used.

Architectural Decisions & Business Value

Decision	Justification	Business Impact
Drop isFlaggedFraud	Prevents data leakage from rule-based system into ML model	Ensures model learns real patterns, not existing rules
PCA at 95% Variance Threshold	Retains maximum signal while removing noise & multicollinearity	Faster training, more stable clusters, reduced overfitting
K-Means on PCA-Reduced Data	Clusters in denoised space produce more meaningful behavioral groups	Richer features for XGBoost; identifies suspicious behavioral profiles
scale_pos_weight in XGBoost	Directly compensates for 1:99 class imbalance without oversampling	Higher Recall — fewer missed frauds; fewer regulatory fines
Recall-Focused Evaluation	A missed fraud is far costlier than a false alarm in AML context	Aligns model optimization with regulatory risk priorities
Incremental XGBoost Update	Adapts to emerging criminal tactics without costly full retraining	Model stays current; no data-retention violations; lower compute cost

Business Value Summary

Reduced False Positives

Behavioral segmentation gives XGBoost contextual signals, reducing noise alerts that overload analyst queues and risk missing the 45-day deadline.

Maximized Fraud Detection

Recall-optimized XGBoost with class imbalance correction catches more real laundering activity, directly reducing regulatory fine exposure.

Adaptive Model Lifecycle

Continuous incremental learning ensures the model evolves with criminal tactics — without expensive full retraining cycles.

Regulatory Auditability

Historical data is never pooled with new batches, maintaining clean data provenance for regulatory review.

This pipeline moves the institution from reactive rule-following to **proactive, adaptive fraud detection** — reducing analyst workload from false positives while maintaining strong sensitivity to genuine laundering behavior. The five-phase architecture (Preprocessing → PCA → K-Means → XGBoost → Continuous Learning) represents a production-ready, regulatory-compliant approach to modern AML at scale.